

# Analyzing the history of *Cognition* using topic models

Uriel Cohen Priva ([Uriel\\_Cohen\\_Priva@Brown.edu](mailto:Uriel_Cohen_Priva@Brown.edu))

Joseph L. Austerweil ([Joseph\\_Austerweil@Brown.edu](mailto:Joseph_Austerweil@Brown.edu))

Brown University, Department of Cognitive, Linguistic, and Psychological Sciences, Providence, RI 02912

## Abstract

There is little to no empirical work analyzing how cognitive science as a field has changed over the last six decades. We explore how *Cognition* changed over the last four decades using topic models. Topic models assume that every word in every document is generated by one of a limited number of topics. Words that are likely to co-occur are likely to be generated by a single topic. We find a number of significant historical trends: the rise of moral cognition, eyetracking methods, and action, the fall of sentence processing, and the stability of development. We introduce the notion of framing topics, which introduce content, rather than present the content itself. These framing topics suggest that over time *Cognition* turned from abstract theorizing to more experimental approaches.

## Introduction

Many researchers are familiar with discussions during post-conference dinners speculating about which research areas are “hot” and lamenting that their own research area has fallen out of favor. Despite these all too common parlor debates, there is little to no empirical work analyzing how cognitive science as a field has changed over the last six decades. Given that many articles have been digitized to be easily accessible electronically and that the abstract and title of most of these articles are free to download, it is now feasible to perform large, data-driven analyses of scientific trends. In this article, we describe one method for analyzing trends in scientific fields, focusing on the journal *Cognition* as part of this special issue.

What is a rigorous, data-driven method for analyzing trends in science? Perhaps the simplest approach would be to count how often particular words and phrases associated with a particular research area are used each year (e.g., “moral”), and analyze any resulting trends. Recently, Behrens, Fox, Laird, and Smith (2012) took this approach for analyzing publication patterns in cognitive neuroscience, and found that particular brain areas were positively correlated with high-impact journals (e.g., the fusiform gyrus). However, this word-based approach is limited by the choice of particular words to model a topic, and runs the risk that changes in language use do not reflect academic interest itself, e.g. artificial intelligence still generates a lot of interest, but it is no longer referred to as artificial intelligence.<sup>1</sup> Rather than use words, Hall, Jurafsky and Manning (2008) demonstrated that *topic models* (Blei, Ng, & Jordan, 2003; Griffiths, Steyvers, & Tenenbaum, 2007) present an appealing approach to tracking the rise and fall of specific research interests in general. Topic models allow us to see how topics, rather than words, trend over time. We followed their approach by analyzing the titles and abstracts of 34 years of *Cognition* articles to track trends in research topics. We

---

<sup>1</sup> As observed in the Google Books Ngram Viewer <http://goo.gl/PUMKQP>

discuss how topics such as *morality* surged, while others such as *developmental* remained stable. Additionally, we extend the work of Hall et al. (2008) in two ways. First, we propose a method for selecting which topics represent consistent trends, regardless of the number of topics used to build the model. Second, we show that interest may lie in what we label *framing topics*, topics that do not model any research domain, but contain words used to frame more contentful topics. Cognition’s abstracts and titles support two framing topics: theory-centric and experimental.

## Topic Models

Although there is no agreed upon representation of word meaning, topic models (Griffiths et al., 2007) provide a relatively simple and practical method for exploring hypotheses about the meanings of words in documents. Topic models assume that every word in every document is generated by one of a number of *topics*. Topics are (Dirichlet-distributed) mixtures of words (so, a topic specifies the probability of a word being produced by that topic), and documents are (Dirichlet-distributed) mixtures of topics. As illustrated in Figure 1, the model generates documents according to a hierarchical process. First, a mixture of topics (the *gist* of the document) is sampled from a Dirichlet distribution. Subsequently each word is sampled from the topics of the document. Documents are biased to be more likely to generate some topics rather than others, and topics are biased to be more likely to generate some words rather than others. Together, these biases lead the model when given a corpus of documents to converge<sup>2</sup> on solutions in which words that are likely to co-occur are generated by the same topic. For example, the “developmental” topic from our analyses assigns high probability to development-related words (such as “child”, “adult”, “young”, and “age” -- see Figure 1d). There is no explicit bias against having a single word generated by more than one topic, and indeed words like “object” in our models are likely to be generated by both a developmental-related topic, and by an eyetracking-related topic.

Our main motivation for using Topic Models is pragmatic. They have been successfully used to analyze trends in other fields of science (notably trends in computational linguistics; Hall et al., 2008). We decided against using a related alternative to Topic Models, Dynamic Topic Models (Blei & Lafferty, 2006), which introduce additional biases regarding the development of topics over time (e.g., trends should change smoothly from year to year). The downside to constraining the development of topics is that the modeling results are more difficult to interpret. How would we know whether a trend found by the model is actually driven by trends in the data or due to the additionally imposed constraints? We also preferred to use Topic Models over another related model, Latent Semantic Analysis (LSA; Landauer & Dumais, 1997), which performs dimensionality reduction on a set of documents to extract a spatial representation of word meanings. The advantage of the topic model approach is in explicitly biasing topics to generate fewer words at high probability, and biasing documents to generate fewer topics at high probability. The topics inferred by the model are therefore more

---

<sup>2</sup> This is done using standard machine learning techniques, see Griffiths, Steyvers, & Tenenbaum, 2007 for more details.

structured, and allow them to explain qualitative and quantitative aspects of word associations better than LSA (see Griffiths et al., 2007, for more details).

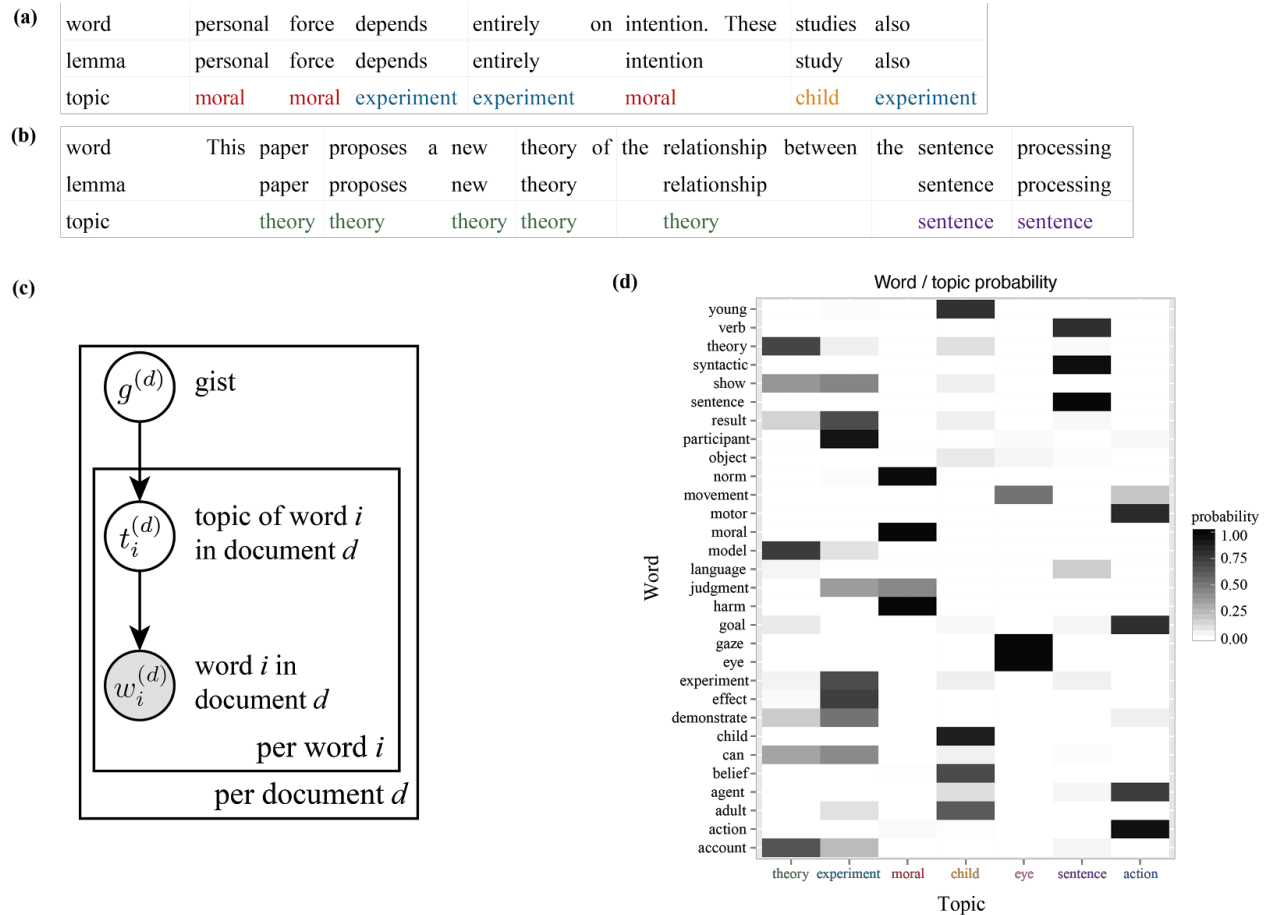


Figure 1: Modeling word meaning using Topic models. (a-b) Nine and twelve word excerpts from Greene et al (2009) and Gibson (1998), respectively. The first row is the original text. The second row is the text after being processed through the lemmatizer. The third row is the topic assigned to each word by the model. (c) A graphical model representation of the Topic model. Each node is a variable in the model and edges encode probabilistic dependencies between the two variables. Each rectangle is called a “plate”, where the variables and edges in the rectangle are copied multiple times (e.g., the inner rectangle represents each word in the current document and each word’s topic). A shaded node encodes that the variable was observed. (d) A word-topic matrix from our trained Topic model limited to seven topics and their most probable words. Note how the *gaze* topic focuses on words related to eyetracking methods and that some words (e.g., movement) are probable in multiple topics (e.g., *eye* and *action*), capturing different senses of the same word.

## Analyzing Trends in *Cognition*

### Materials and Methods

We retrieved<sup>3</sup> from ScienceDirect all the titles and abstracts of *Cognition* articles published between 1980 and 2014 (3104 abstracts), and loaded them to R (R Core Team, 2014). We cleaned the data by removing all punctuation, single-character words, HTML escape codes, English function words (e.g. *this*, *being*; using a stopwords list from the *tm* library, Feinerer and Hornik, 2014), collapsing uppercase and lowercase distinctions, and converting all words to their lemmas, using the WordNet-based (Princeton University, 2010) NLTK lemmatizer (Bird et al, 2009). Thus a fragment such as “*LSA, and topic models*” became “*lsa topic model*”. After the cleaning process was done, we excluded abstracts that had no words in the abstract (mostly non-articles such as short responses, indices and contents of volumes). The resulting set of 2512 abstracts had 107 words per abstract on average. We used the R *lda* library (Chang, 2012) to train topic models for multiple numbers of topics (between 30 and 50) using at most 1000 iterations. The hyperparameters,  $\alpha$  and  $\eta$ , were both set to 0.1, which encourages the model to assign topics to documents such that each document is composed of a few topics and to learn topics that produce a few words with high probability, respectively. The main output of each model is an assignment of topics for every word in the dataset. For instance, a word vector of the form “*lsa topic model*” may be translated to “2 2 10”, in which the first two words were generated by topic 2, and the third word by topic 10.

To analyze trends in topics over time, we used the output from the model to calculate the weighted contribution of each topic  $t$  in each year  $y$  using Equation (1). It calculates for each year the proportion of words generated by each topic, but controls for abstract length by dividing each word’s contribution by the abstract’s length. Thus, the yearly topic contribution estimate,  $p_y(t)$ , is defined to be

$$\frac{1}{|D_y|} \sum_{d \in D_y} \frac{|\{w \in d: \text{topic}(w) = t\}|}{|d|} \quad (1)$$

where  $t$  is a topic,  $w$  is a word in document  $d$ , and  $d$  is a document in  $D_y$ , the collection of documents published in year  $y$ .

Because the choice of the number of topics can affect the topic trends, we needed a method for detecting which topic trends are likely to be meaningful and robust. After experimenting with a few methods and thresholds, we decided to include topics that were in the top  $1-2.5/(\text{number of topics})$  percentile of topics (top 10% for 25 topics, top 5% for 50 topics), or topics that had a significant trend ( $R^2 > 0.4$ ) when trying to predict trends in topic contribution over the years, using a cubic-function regression. To prevent any influence of the particular choice of  $K$  (the number of topics learned by the model), we discuss only those topics that passed the above criterion for nearly every  $K$ . The presented results are the most

---

<sup>3</sup> Retrieved 2014-07-28

conservative of the different options that we tried and thus, are most likely to be robust to the choice of  $K$  and thresholds.

## Results and Discussion

There were three types of topics inferred by the model that passed our criterion: *framing* topics, topics with major trends, and topics that are stable, yet still occur relatively often. We discuss each of these in turn and then discuss some trends that were less robust (occurred roughly half of the time).

Framing topics do not concern any particular research domain, but contain words that are used to frame more contentful topics. They arise for the following reason: When using topic models, every word in a document has to be generated by some topic. For topics that are likely to generate words such as *development* and *child*, their interpretation is straightforward—the topic is related to developmental psychology. But some words such as *show* and *demonstrate* are not about any particular research interest. Instead, they are words that frame the content of the abstract. The model still has to generate these words, and it is biased against using the contentful topics to generate them, as every contentful topic will have to allocate some probability mass to content-framing words. Instead, in every model trained with more than 10 topics, one topic emerged that generated the framing context. In models with 30 or more topics, two framing topics emerged rather than one.




Figure 2: Trends in the prevalence *theory-framing* topic (red), *experiment-framing* topic (blue), and the sum of the prevalence two topics (black). Although the prevalence of the two topics remains stable, the *experiment-framing* topic has become more prevalent while *theory-framing* declines. The curves are best-fit polynomial functions using the loess R function and shaded area around the curves denotes 95% confidence intervals.

The two framing topics were the *theory-framing* topic, which generates words related to abstract theories of cognition, and the *experiment-framing* topic, which generates words related to empirical experiments testing aspects of cognition. Figure 2 plots the yearly trend of each framing topic (red for *theory-framing* and blue for *experiment-framing*) along with the sum of the two trends (black). There is a clear downward trend for the theory-framing topic, whereas there is a clear upward trend for the experiment-framing topic. However, the trend of their sum remains stable over time. This suggests that over the last three decades, *Cognition* articles are increasingly framed in terms of experiments, rather than abstract theories. That is, over time the experiment-framing topic is replacing the theory-framing topic.

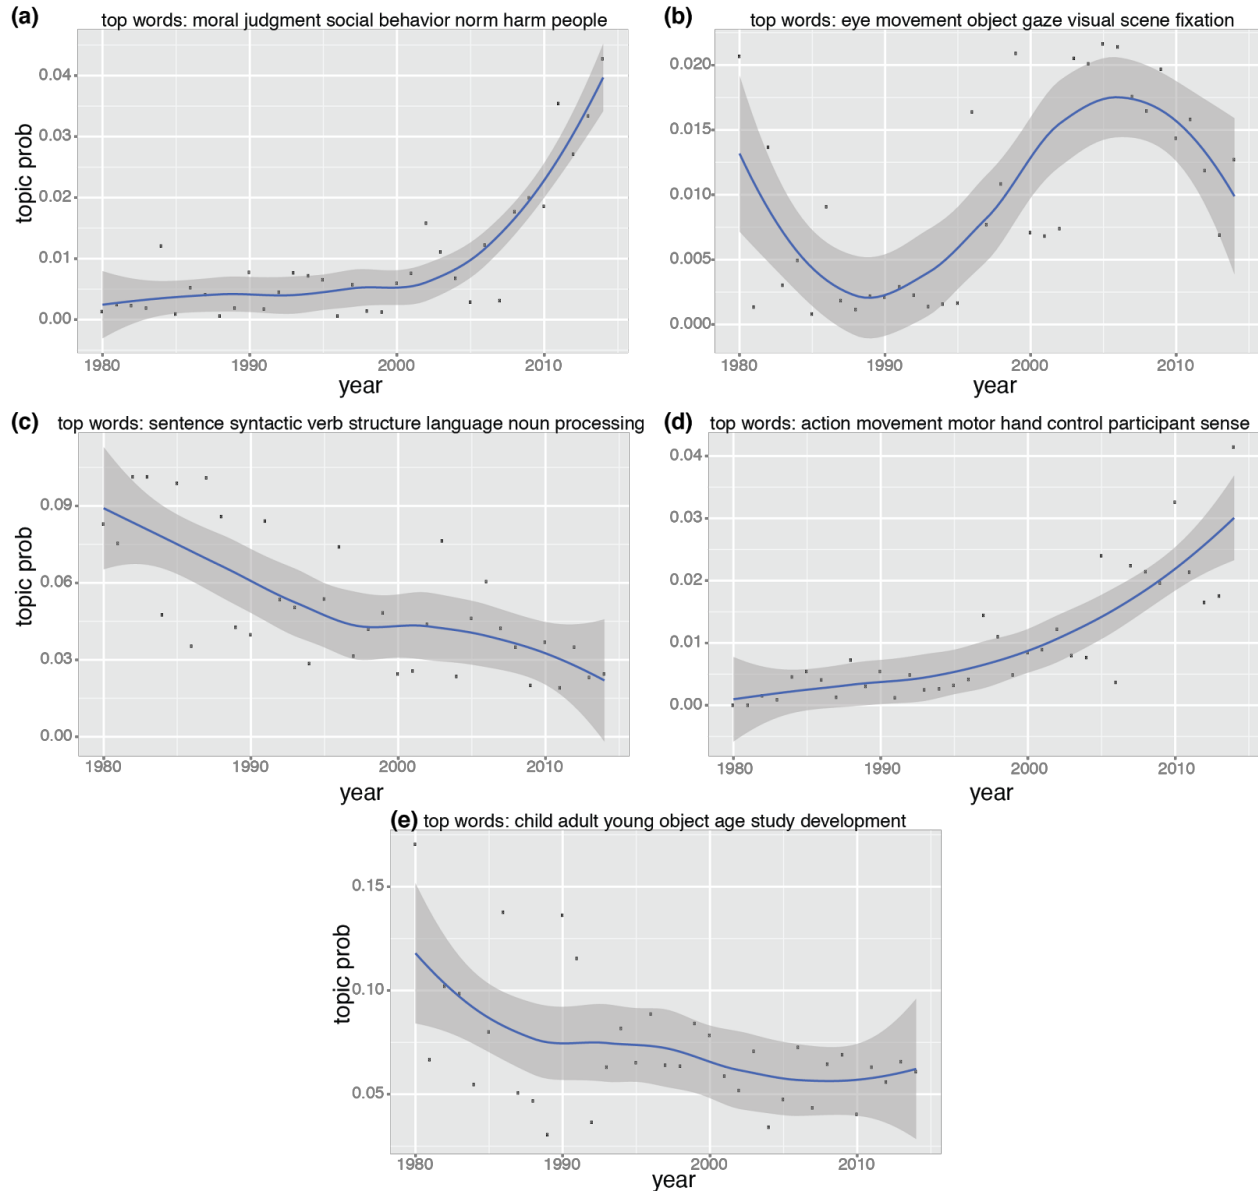


Figure 3: Trends of topic prevalence in *Cognition* passing our criterion for the (a) *moral*, (b) *gaze*, (c) *sentence processing*, (d) *motor-action*, and (e) *development* topics. The curves are best-fit polynomial functions using the loess R function and shaded area around the curves denotes 95% confidence intervals. See text for more details.

As shown in Figure 3, we uncovered five topics that passed our criterion for nearly all  $K$ , a *moral* topic covering moral, social, and emotional aspects of cognition, a *gaze* topic covering eye tracking methodologies outside of vision (e.g., sentence processing or moral cognition), a *sentence processing* topic, a *motor-action* topic, and a *development* topic. Although the *moral* topic has very low probability in *Cognition* during the 1980s and 1990s, its prominence has skyrocketed in the last decade from producing around 0.5% of the words in the mid-2000s to producing about 4% in 2014. This is a few years after groundbreaking work by Greene and

Haidt. At the time of writing this article Greene et al 2001, and Haidt 2001 are widely cited (over 2000 citations each). The trend of the *gaze* topic is more complex. The most striking feature is the sharp increase in its popularity from nearly 0% at the early 1990s to almost 2% in the mid 2000s, which then plateaued (and possibly now is in a downward trend). The sharp increase in the 1990s seems to be due to the sudden popularity of using eye tracking as a method for analyzing language processing, which coincides with Tanenhaus, Spivey-Knowlton, Eberhard, and Sedivy (1995)'s seminal article. The sentence processing topic has been trending approximately linearly down in *Cognition* since the beginning of the 1980s, from nearly 9% to close to 3%. Conversely, motor action trends positively, from nearly 0% to about 3%. Although the *development* topic also passes our criterion, unlike the other discussed topics there are no notable trends. It has been generating approximately the same proportion of topics over the last few decades.

## Summary and Concluding Remarks

In this paper we used topic models to explore how different topics have trended over the history of *Cognition*. While one may expect specific topics would rise or fall, changes in framing topics reveal tendencies in the journal or field as a whole. One of the most robust trends in *Cognition* is the rise of the experimental approach at the expense of more abstract theorizing. We believe that topic models present a methodical way to analyze how scientific interest changes over time.

## Acknowledgements

Both authors contributed equally to this paper. This paper greatly benefited from discussion with Steven Sloman. We thank Scott AnderBois, Elinor Amit, and Russell Church for their feedback on earlier versions of the manuscript.

## References

- Behrens, T. E. J., Fox, P., Laird, A., & Smith, S. M. (2013). What is the most interesting part of the brain?. *Trends in Cognitive Sciences*, 17(1), 2-4.
- Bird, S., Loper, E., & Klein, E. (2009), *Natural Language Processing with Python*. O'Reilly Media Inc.
- Blei, D., & Lafferty, J. D. (2006). Dynamic topic models. In *Proceedings of the 23rd International Conference on Machine Learning* (pp. 113-120). ACM.
- Blei, D. M., Ng, A. Y., & Jordan, M. I. (2003). Latent Dirichlet allocation. *Journal of Machine Learning Research*, 3, 993-1022.
- Chang, J. (2012). *lda*: Collapsed Gibbs sampling methods for topic models. R package version 1.3.2. <http://CRAN.R-project.org/package=lda>.



Feinerer, I. and Hornik, K. (2014). tm: Text Mining Package. R package version 0.5-10.  
<http://CRAN.R-project.org/package=tm>.

Gibson, E. (1998). Linguistic complexity: Locality of syntactic dependencies. *Cognition*. Vol 68, issue 1, 1-76.

Greene, J.D., Cushman, F.A., Stewart, L.E., Lowenberg, K, Nystrom, L.E & Cohen, J.D. (2009). Pushing moral buttons: The interaction between personal force and intention in moral judgment. *Cognition*. Volume 111, issue 3, 364-371.

Greene, J.D., Sommerville, R.B., Nystrom, L.E., Darley, J.M., & Cohen, J.D. (2001). An fMRI investigation of emotional engagement in moral Judgment. *Science*, Vol. 293, 2105-2108.

Haidt, J. (2001). The emotional dog and its rational tail: A social intuitionist approach to moral judgment. *Psychological Review*. 108, 814-834.

Griffiths, T. L., Steyvers, M., & Tenenbaum, J. B. (2007). Topics in Semantic Representation. *Psychological Review*, 114(2), 211-244.

Hall, D., Jurafsky, D. & Manning, C. D. (2008). Studying the history of ideas using topic models. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing* (pp. 363-371). Association for Computational Linguistics.

Landauer, T. K., & Dumais, S. T. (1997). A solution to Plato's problem: The latent semantic analysis theory of acquisition, induction, and representation of knowledge. *Psychological Review*, 104, 211-240.

Princeton University (2010). "About WordNet." WordNet. Princeton University. 2010.  
<<http://wordnet.princeton.edu>>

R Core Team (2014). R: A language and environment for statistical computing. R Foundation for Statistical Computing, Vienna, Austria. URL <http://www.R-project.org/>.

Tanenhaus, M. K., Spivey-Knowlton, M. J., Eberhard, K. M., & Sedivy, J. C. (1995). Integration of visual and linguistic information in spoken language comprehension. *Science*, 268(5217), 1632-1634.